

Predictive AI for Engineers — Detailed Lesson Plan

Text & Speech Analytics · four hands-on sessions of 90 minutes · a no-code lab for HAL engineers

This is the working lesson plan for a one-day, four-session introduction to **predictive (classical) AI** through the two kinds of messy data engineers deal with every day: written **text** and spoken **speech**. It is the detailed companion to the short coverage note in `SESSION_PLAN.md` — here you get the per-session timings, the exact demos, the case studies (with honest caveats and sources), the discussion prompts, and how we handle the skeptics.

The whole day deliberately covers the **predictive / analytical** half of AI — machine learning that *learns a pattern from past examples and predicts the next case*. Generative AI (ChatGPT and its relatives) is a different tool for a different job and is covered separately by other speakers. Saying that out loud, early and often, is part of the plan.

Who this is for, and the one goal

The room is **experienced mechanical, aeronautical, production and quality engineers** with no programming background — people who know their machines far better than we know them. Most arrive holding one of two beliefs: **"AI can do anything"** or **"AI is marketing nonsense."** Both are wrong, and the truth is more useful than either.

Our single goal: every person leaves able to **reason about an AI claim** instead of believing or dismissing it — *what did it learn from, how was it tested on unseen cases, how confident is it, and where does it get things wrong?* We earn that by showing real, **honest, imperfect** results on data that looks like theirs. One overhyped claim and the skeptics tune out for the rest of the day, so credibility is the currency we spend carefully.

Format — a no-code lab

- **No coding required.** Every code cell is pre-written. Participants only **run** cells (Shift + Enter) and read the output; in each session they change **one highlighted value** and re-run to see the result move.
- Runs in **Google Colab** in a browser — nothing to install. The text and audio notebooks are self-contained (built-in data, fixed seeds, so everyone sees the same numbers); one notebook fetches a small public dataset live to prove the methods hold on real data.
- A rhythm repeated every session: **short framing** → **run a cell** → **"what just happened"** → **a plain-English vocabulary card** → **a knob to turn** → **recap**.
- **Mathematics is shown visually only** — a confusion-matrix grid, a similarity heatmap, a fitted line, a feature-importance bar, a spectrogram. Techniques are named for honesty (TF-IDF, logistic regression, FFT) but **no equations and no recall are expected**.

The day at a glance

#	Session (90 min)	Notebook	The one new idea	The money-shot
1	Text I — sorting the words you already write	<code>04_text_analytics.ipynb</code>	A computer can't read words, only count them: turn words into numbers (TF-IDF), then it's ordinary classification	Top clue-words per team — the model's reasoning in plain English
2	Text II — predict, extract & find repeat faults	<code>04b_text_predict_extract.ipynb</code>	The same words-to-numbers front end answers many questions — a category, a number, a duplicate, a part number	A similarity heatmap lighting up the same chronic fault, worded five different ways
3	Speech I — voice → text → decision	<code>05_speech_analytics.ipynb</code>	Speech-to-text is "audio in, text out"; then it's the text analytics from Sessions 1–2	A low word-error transcript that still gets the one part number wrong
4	Speech II — predicting from the sound itself	<code>sound_fault_prediction.ipynb</code>	A sound is already numbers; learn the pattern and a computer can <i>hear</i> a fault — no words at all	Three fault types separating into clean clusters; the model leaning on the same whine and rumble a technician hears

Suggested clock: 09:30 S1 · 11:00 S2 · (lunch) · 13:30 S3 · 15:00 S4 · ~16:30 close. Breaks between sessions; all timings flexible.

The single thread (put this on a wall poster)

Every session — text, speech, sound — is the *same five steps*. Master the shape once and the rest is variations:

dataset → split (hide some) → train → predict → measure (honestly).

Keep pointing back to it. The data changes shape (sentences, audio, a spectrum of pitches) but the picture never does. By Session 4 the room should be finishing the sentence for you.

How we win the room (the doctrine behind every session)

- **Lead with credibility, not capability.** Open by naming the two mindsets and promising the honest middle. The first example of the day is the everyday **email spam filter** — classical text AI (bag-of-words, decades old, deployed at massive scale) that nobody calls magic. (*Note: attribute it to spam filtering in general, not to any one provider's current algorithm.*)
- **Always show an imperfect result and open the confusion matrix.** For experienced engineers, visible errors are a *trust* signal; a flawless demo reads as a vendor trick. Every "it can" is paired with an honest "it can't."
- **The accuracy paradox, in their language.** If real failures are 1-in-1000, a model that always says "healthy" is 99.9% accurate and useless. That is why precision/recall and the confusion matrix — not a single accuracy number — are what an engineer should demand.
- **Predict-then-reveal, one knob.** Have the room commit to a guess out loud, change one highlighted value, re-run. Active prediction plus a single controlled variable produces the "aha" that passive watching never does.
- **Show, don't state.** Charts, never equations. Engineers read Pareto and control charts fluently; one formula on screen and a no-code audience checks out.
- **Respect their expertise.** Frame every text technique as *"reading the words you already write"* and every audio technique as *"a crude version of what your ear already does."* The model just learns their existing labelling habits; it rediscovers the cues they already trust; every extraction is auditable. **Auditability is the selling point.**
- **One aerospace + one factory example per technique**, so the skill is seen as general, not an aviation trick.
- **Be explicit about real vs synthetic vs simulated** — this audience will ask. Workshop demo data is synthetic and self-contained; named datasets (ASRS, CWRU, MIMII, MAFAULDA) are real; C-MAPSS is physics-simulated.
- **Close every demo with a prompt that pulls THEIR data into the room.** "What is the label column in *your* snag logs? Who assigns it today? Where would a wrong prediction actually cost you?" Concrete prompts get debate from skeptics; "any questions?" gets silence.

Vocabulary, spread across the day

We do **not** front-load a glossary. Each term is introduced the moment it is first needed, then reused relentlessly.

Term	First met	In one plain line
dataset, row, column	S1	a table — one row per case, columns for what we recorded
feature / label / label column	S1	features = the clues we get; label = the answer we want; the label column holds the answers
train/test split	S1	learn from most, hide some to test honestly on unseen cases
model, accuracy	S1	the learned pattern; the fraction right on the hidden set
bag-of-words, TF-IDF, stop words	S1	counting words; rare distinctive words count for more; ignore "the/and"
confusion matrix	S2	a grid of what it got right and <i>how</i> it got things wrong
classification vs regression	S2	predict a bucket vs predict a number
cosine similarity	S2	0–1 score for how much two reports use the same distinctive words
entity extraction (NER)	S2	pulling part numbers / components out of free text
topic modelling (LDA)	S2	grouping co-occurring words into themes with no labels
data leakage	S2	giveaway clues sneaking into the test — a too-perfect score
ASR, acoustic/language model	S3	speech-to-text; the two halves that turn sound into likely words
word error rate (WER)	S3	the speech version of accuracy — wrong + missing + extra words
keyword spotting	S3	listening for a few fixed commands instead of every word
sample rate, FFT, spectrum	S4	sound is numbers; the FFT splits it into the pitches it's made of
spectrogram, MFCC	S4	a picture of pitch over time; the "professional" sound features
supervised vs anomaly detection	S4	learn from labelled faults vs learn "normal" and flag the surprising
orders (1×, 2×)	S4	fault pitches read as multiples of shaft speed, not raw Hz

Session 1 — Text I: sorting the words you already write

Notebook: 04_text_analytics.ipynb · **New idea:** words become numbers, then it's ordinary classification.

By the end, participants can: name the five steps of a predictive-AI project; use *dataset* / *feature* / *label* / *train-test split* / *accuracy* correctly; explain (in plain words) how text becomes numbers; and read a model's top clue-words to see it is not a black box.

Time	Segment	What happens, and how it lands
0:00–0:08	Frame the day	Name the two mindsets. Promise the honest middle. The spam-filter anchor: you already trust classical text AI. State plainly: this is <i>predictive</i> , not generative.
0:08–0:18	The five-step poster + first words	Walk the spine: dataset → split → train → predict → measure. Introduce dataset/feature/label/label column as a plain spreadsheet idea.
0:18–0:23	Into Colab	Orient: Shift+Enter, run top-to-bottom, "Restart and run all" fixes almost everything.
0:23–0:53	Run Notebook 04	Words → numbers (bag-of-words, stop words, TF-IDF) shown on two sentences first; vectorize 48 reports; train/test split; honest accuracy; top clue-words per team (oil/pressure/fluid → Hydraulic). Predict four fresh reports incl. one deliberately ambiguous — watch confidence drop.
0:53–1:00	Turn the knob	Change the train/test split ratio and re-run; discuss why the score wobbles on a tiny set.
1:00–1:12	Where this lives at work	Aerospace: auto-routing snags by ATA chapter; tagging ASRS-style narratives. Industrial: CMMS/IT ticket triage (an SVM hit ~85% on 3,632 HVAC maintenance requests); spam filtering.
1:12–1:22	Discussion	"What is the label column in <i>your</i> logs? Who assigns it today, and how long does it take?"
1:22–1:30	Recap + bridge	Five steps revisited; tease S2: sorting was the doorway — next we predict urgency, repair time, and catch repeat faults.

Skeptic handling. *"This is just if-then rules / statistics."* — Yes, and that's the point: predictive AI is disciplined pattern-finding. The top-clue-words list proves it learned something human-readable, not magic. *"Why isn't it 100%?"* — On 48 toy reports the score is a rough indicator; a model that claims 100% on real data usually peeked at the answers (we prove this in S2).

Real, citable anchors: scikit-learn's own text tutorial reports ~83% (Naive Bayes) and ~91% (linear model) on a real 4-category corpus — honest mid-80s/low-90s numbers to set expectations. NASA's ASRS holds 2.3M+ de-identified narratives (free, public-domain) that look just like HAL snag/incident logs; NASA itself calls them "soft data" it does not verify — an honest data-quality lesson in one sentence.

Session 2 — Text II: predict, extract & find repeat faults

Notebook: `04b_text_predict_extract.ipynb` (new) · **New idea:** the same TF-IDF front end answers many questions — just swap the final step.

By the end, participants can: tell classification from regression and know which question each answers; read a confusion matrix for *what* a model gets wrong; explain how duplicate/recurring faults are found with no labels; see how part numbers are lifted out of prose and audited; and recognise data leakage (the too-perfect score).

Time	Segment	What happens, and how it lands
0:00– 0:06	Recall + preview	One-line recap of "words → numbers." Today: predict a category, predict a number, find repeats, pull out parts.
0:06– 0:12	The snag log	Build a 300-row synthetic MRO log: free-text report, system, urgency, downtime hours. Note the last is a <i>number</i> .
0:12– 0:28	Predict a CATEGORY (urgency)	TF-IDF + logistic regression → ~81% on unseen snags. Open the confusion matrix : mistakes are mostly Low↔Medium; it caught 10/15 High snags and slipped High→Low only once — <i>rare, but it happened</i> , which is exactly why the High bucket keeps a human's eyes.
0:28– 0:42	Predict a NUMBER (downtime)	Same TF-IDF in, swap to a regressor → off by ~1.5 h on ~6 h jobs ($R^2 \approx 0.6$). The predicted-vs-actual scatter makes "it learned the relationship" visible — and honest about the big-job misses. <i>This is the classification-vs-regression "aha"</i> .
0:42– 0:56	Find RECURRING faults	TF-IDF + cosine similarity on snags worded differently. The heatmap lights up two blocks — the same chronic fault, de-duplicated, no labels. The single most HAL-relevant text demo.
0:56– 1:06	Pull out PART NUMBERS	A few readable rules extract part numbers into an auditable table, then a "most-mentioned parts" chart — free text becomes a reliability/spares signal. Every hit is checkable.
1:06– 1:16	Bonus: themes (LDA) + the honest-test trap	One pre-baked LDA topic view (some topics map to systems, some to severity language — <i>you</i> name them). Then the same pipeline on a real public corpus: ~84% honestly vs ~95% when giveaway headers leak in. <i>How was it tested, and could it have peeked?</i>
1:16– 1:30	Discussion + recap	"Where would a wrong prediction be expensive vs cheap for you? Do your historical logs exist <i>labelled</i> ?"

Skeptic handling. *"Can it replace our engineers?"* — No: it hands you a prioritised list and a reason; a human decides. The confusion matrix and the High→Low slip make the limits measurable, which is exactly why it's trustworthy.

Real, citable anchors (with honest caveats):

- **Chronic/recurring-defect detection** is a live commercial product (Veron's ChronicX). Its figures — used by >25% of the world fleet, ~40% of defects unflagged due to ATA mis-coding, 90%+ ATA prediction — are **vendor-claimed, not independently audited**. Lead with the capability (de-duplicating differently-worded faults), not the percentages.
- **Part-number extraction**: a US DoD/AFIT study lifted part-number extraction F1 from ~3% (plain regex) to ~95% (a trained NER model) — an honest *mixed* result (rigid codes like NSNs were easy for both).
- **Defence precedent**: a national defence organisation (DSTG/RMIT, ICAS 2024) mined two years of military-aircraft maintenance + pilot free-text to surface "lead indicators" of defects — the closest published analogue to HAL's situation.
- **Topic modelling, honestly**: a 2021 NASA-data study recovered >70% of hand-curated COVID-19 ASRS reports via topic modelling — *but* a separate study concluded off-the-shelf topic tools were "insufficient for sense-making" on their own. Useful map, not a verdict.
- **Industrial transfer**: the NHTSA consumer-complaints database (~1.9M free-text records, public) is the warranty-mining analogue.

Indian context to mention lightly: after a 2025 accident, India's DGCA is reported to be tightening repetitive-defect reporting (a defect recurring three times within 15 flight cycles) — **press-reported/draft**, so cite it as motivation, not settled regulation.

Session 3 — Speech I: voice → text → decision

Notebook: `05_speech_analytics.ipynb` · **New idea**: speech-to-text is "audio in, text out"; then it's the text analytics you already know.

By the end, participants can: describe ASR as the same predict-from-examples idea with sound as the input; read WER as the speech version of accuracy and explain why a low WER can still be unsafe; and say why a small fixed vocabulary (keyword spotting) is more reliable than full transcription.

Time	Segment	What happens, and how it lands
0:00– 0:08	Frame	The last messy data type: spoken words. The pipeline: speak → transcribe → decide. Draw the classical→modern line (HMM-GMM → deep models, ~2012) explicitly so nobody thinks "this is ChatGPT" — both are <i>predictive</i> , not generative.
0:08– 0:30	Run Notebook 05 (pipeline)	Manufacture a spoken report (text-to-speech, no mic needed) → transcribe with Whisper → route with the Session-1 classifier, with a confidence bar. (<i>First cell installs two tools, ~1 min — the only setup of the day.</i>)
0:30– 0:45	Measure it honestly (WER)	A tiny self-contained WER demo: a noisy transcript scores a <i>worse</i> 15% yet stays harmless; a "cleaner" 8% transcript quietly turns MS21042 into MS21052 — the one word that matters, now wrong. Low error rate ≠ safe . Rule: a human verifies numbers and serials.
0:45– 0:55	Keyword spotting	Listening for a 4-command vocabulary inside messy transcripts — robust to stutters and filler, and it leaves the no-command line alone. Why hands-free shop-floor tools use wake-words, not dictation.
0:55– 1:08	Where it works, where it breaks	Hard by default (accents, jargon, radio noise), strong with domain data: the EU HAAWAI ATC project reached >95% (controllers) / >90% (pilots) word recognition <i>after</i> domain training, ~97% on callsigns when fused with radar. Open ATC corpora: ATCOSIM (clean), ATCO2 (noisy).
1:08– 1:20	Discussion	"Where do people <i>speak</i> information that never gets written down — handovers, inspections, calls? What's the one word in your world that must never be mis-heard?"
1:20– 1:30	Recap + bridge	Speech → text → decision, honestly measured. Tease S4: next we throw the words away and predict from the <i>sound itself</i> .

Skeptic handling. *"ASR is hype — it never gets my accent."* — Agree out loud, then show the WER demo and the HAAWAI counter-evidence: out of the box it struggles; trained on your domain it gets strong. Honesty about the limit is what makes the capability believable.

Honest note on Whisper: trained on 680,000 hours of audio; it is excellent zero-shot but does **not** beat specialist models on their home turf and can occasionally hallucinate text — one more reason critical fields are human-verified.

Session 4 — Speech II: predicting from the sound itself

Notebook: `sound_fault_prediction.ipynb` · **New idea:** a sound is already numbers; learn the pattern and a computer can *hear* a fault — no words.

By the end, participants can: explain that a sound is just numbers and that a few features (energy, low/high bands, dominant pitch) capture what the ear hears; tell supervised classification from unsupervised anomaly detection and which fits the realistic "faults are rare" case; and read fault pitches as orders of shaft speed.

Time	Segment	What happens, and how it lands
0:00–0:08	The technician's ear	A seasoned tech walks past and says "that bearing's going" — from the sound. Can a computer learn that? Same idea, last time: numbers → pattern → judge.
0:08–0:22	Sound is numbers; listen, then featurise	Sample rate → a clip is a list of numbers. <i>Play</i> healthy / bearing / imbalance clips. The FFT splits a sound into its pitches (a prism); five features capture loudness, low rumble, high whine, dominant pitch, brightness.
0:22–0:34	See the pattern, then classify	Scatter of low-band vs high-band energy → three clean clusters. Train/test split → ~88% on unseen clips. Feature importance shows the model leaning on the whine and the rumble — the very cues the ear uses.
0:34–0:44	Diagnose a brand-new machine	Play a fresh clip, get a fault call + confidence. Turn the knob: try other faults / seeds; widen the bearing range and watch accuracy soften — honest.
0:44–0:58	When you have no fault examples	The realistic aerospace case: lots of healthy sound, few faults. Train an anomaly detector on healthy only (AUC ≈ 0.94); the histogram shows healthy left, faults right of the alarm line. Move the line: ~5% false alarms catches ~96% of imbalance and ~ $\frac{2}{3}$ of bearing faults — <i>no free lunch</i> , and the trade-off is visible.
0:58–1:10	Real datasets + orders	Name the real, benchmarked field: CWRU bearings, MAFAULDA (with a mic channel), MIMII / the DCASE challenge (honest AUCs ~0.54–0.96). Read pitches as orders (imbalance 1×, misalignment 2×), not raw Hz.
1:10–1:20	Where it lives in aerospace	Helicopter HUMS detects ~70% of the failure modes it monitors (strong, explicitly imperfect); engine EHM (Rolls-Royce monitors 8,000+ aircraft); acoustic-emission NDT for cracks — same recipe: signal → features → classify/flag.
1:20–1:30	The day in one picture + close	Numbers, text, speech, sound — it was always the same five steps. The "neither magic nor nonsense" close: you can now <i>question</i> an AI claim. Hand off to the generative-AI sessions to come.

Skeptic handling. *"But we don't have failure labels."* — Exactly why the anomaly-detection demo matters: learn normal, flag the surprising, no fault examples needed. *"Why not 100%?"* — DCASE's public scores run from ~54% to ~96%; a flawless demo is the thing to distrust.

Case studies you can cite (and how to cite them honestly)

Case	What it really is	How to say it
Airbus Skywise — A330neo bleed valve (EASA emergency AD, Aug 2022)	Cross-fleet data analysis flagged a software bug over-stressing a valve pin; caught <i>before</i> any accident	The best honest safety story for the room — predictive analytics as an engineering early-warning, not cost-savings hype. (Skywise now connects 12,000+ aircraft — cite the live figure.)
Delta TechOps on Skywise	Maintenance-caused cancellations fell from 5,600+ (2010) to ~55 (2018), >95% success on pending-failure predictions	Powerful — but Delta self-reported , not independently audited. Model the "how was this measured?" discipline.
Veryon ChronicX	NLP that recognises the same chronic defect described differently and predicts ATA codes	Vendor-claimed figures; lead with the capability, flag the numbers as unaudited.
EU HAAWAI	Domain-adapted ATC speech recognition	>95% / >90% word recognition <i>after</i> domain training; ~97% callsigns with radar. (Say "improved roughly twofold with an hour of training data"; the exact "50%→25% WER" phrasing is not in the source.)
DCASE Task 2 (MIMII/ToyADMOS)	Public benchmark for detecting machine faults from sound using only normal data	The anti-hype anchor: honest AUCs ~0.54–0.96 across machine types.
Rolls-Royce EHM / TotalCare	Engine health monitoring underpinning "power-by-the-hour"	"Monitors 8,000+ aircraft" (vendor figure; the page says aircraft). Shows the business model that forces honest models.
IAF–IIT Bombay, Su-30 MKI (2026)	Indigenous prognostic maintenance for the ~270-strong fleet	"This is happening here" context for HAL — but it is sensor/prognostics , not text/audio; cite as context only.

On HAL itself: there is **no verified public evidence** of HAL running predictive ML in production. Frame everything as "**applicable to HAL's operations**," not "HAL already does this." HAL's strongest latent asset is its own historical snag/overhaul logs — *if* they are clean and labelled.

Public datasets (Colab-usable or citable)

Dataset	What	Use	Where
NASA ASRS	2.3M+ de-identified aviation safety narratives (free text + synopsis)	S1–S2 text routing / topics	asrs.arc.nasa.gov (public domain; 10k/download)
MaintNet	6,169 real aviation maintenance logbook entries + tools	S1–S2 "cleaning beats the model"	aclanthology.org/2020.coling-demos.2
FAA SDRS	~1.7M service-difficulty records, free-text + structured	S1–S2 trend/routing	faa.gov/av-info/download_SDR
20 Newsgroups	~18k real docs, one-line fetch	S2 real-data proof + header-trap	built into scikit-learn
ATCOSIM / ATCO2	Clean / noisy ATC speech corpora	S3 realism	spsc.tugraz.at · arxiv.org/abs/2211.04054
Speech Commands	1-sec spoken-word clips	S3 keyword spotting	arxiv.org/abs/1804.03209
CWRU Bearing	The bearing-fault benchmark (.mat)	S4 "real, benchmarked field"	engineering.case.edu/bearingdatacenter
MIMII	Real valve/pump/fan/slide-rail sound + factory noise	S4 anomaly detection	zenodo.org/records/3384388
MAFAULDA	1,951 runs: imbalance/misalignment/bearing, incl. a mic channel	S4 maps to hand-diagnosed faults	www02.smt.ufrj.br/~offshore/mfs
NASA C-MAPSS	<i>Simulated</i> turbofan run-to-failure (RUL)	S4 regression bridge (label it simulated)	data.nasa.gov
AI4I 2020	10k-row synthetic predictive-maintenance table	accuracy-paradox teaching	archive.ics.uci.edu/dataset/601

Keep heavy downloads (CWRU, MIMII, MAFAULDA) out of the live lab — name them and link them; the synthetic demos are the floor, not the ceiling.

Lab logistics & pre-flight

- A computer lab with **internet**, one machine per participant (pairs are fine), able to reach **Google Colab** and `pypi.org`.
- A **Google account** per participant (personal works), or use the upload-the-notebook fallback.
- A projector for the framing slides and the money-shot charts.
- **Pre-flight (day before):** run **Session 3** end-to-end once — its first cell `!pip install -q gTTS openai-whisper` takes ~1 minute and downloads a small model on first transcription; confirm the network allows it. Everything else needs no install. Keep one pre-run copy of each notebook open as a backup.
- **The one reset that fixes almost everything:** *Runtime* → *Restart and run all*. Say it often.

Mapping to the existing material

- **Sessions 1–4** use `04_text_analytics`, `04b_text_predict_extract` (new), `05_speech_analytics`, and `sound_fault_prediction`.
- The **numbers-first notebooks** `00` – `03` (intro vocabulary, classification + the accuracy trap, regression & forecasting, predictive maintenance + feature importance + anomaly detection) are an excellent **optional pre-read or appendix** — they teach the same five steps on sensor numbers and are where the *accuracy trap* and *feature importance* ideas were first built.
- Slides, speaker notes and the take-home field guide already exist and support the framing throughout.

Sources & further reading (selected, verified)

Topic	Source
Spam filtering as classical text AI	en.wikipedia.org/wiki/Naive_Bayes_spam_filtering
Honest text-classifier accuracy + header trap	scikit-learn.org (20 Newsgroups tutorial & dataset notes)
Work-order text classification (SVM ~85% / 3,632 HVAC)	ASCE J. Arch. Eng. (2022), doi 10.1061/(ASCE)AE.1943-5568.0000522
ASRS database & "soft data" caveat	asrs.arc.nasa.gov
Part-number NER (3%→95%)	apps.dtic.mil/sti/trecms/pdf/AD1157022.pdf
Defence maintenance-text NLP (DSTG/RMIT, ICAS 2024)	icas.org/icas_archive/icas2024 (#0083)
Topic modelling on ASRS (>70% recovery; "insufficient for sense-making")	ntrs.nasa.gov/citations/20210011508 ; /20205003750
WER definition	en.wikipedia.org/wiki/Word_error_rate
HAAWAI ATC ASR	cordis.europa.eu/article/id/442201
Deep models replace GMM (~2012)	cs.toronto.edu/~hinton/absps/DNN-2012-proof.pdf
Whisper (680k hours; not SOTA on LibriSpeech; hallucinations)	arxiv.org/abs/2212.04356
DCASE 2020 Task 2 (AUC ~0.54–0.96)	ar5iv.labs.arxiv.org/abs/2006.05822
Helicopter HUMS (~70% of monitored failure modes)	skybrary.aero/articles/vibration-health-monitoring-vhm
Skywise A330neo early warning (EASA AD 2022)	theaircurrent.com/aviation-safety/skywise-big-data-airbus-...
Delta TechOps figures (self-reported)	deltatechops.com
Rolls-Royce EHM	rolls-royce.com/products-and-services/civil-aerospace/services/ehm.aspx
IAF–IIT Bombay Su-30 MKI PdM	indianmasterminds.com (defence)

All quantitative claims above were fact-checked; figures that are vendor-claimed, self-reported, simulated, or press-reported/draft are flagged as such in the text. Use them in that spirit — modelling the exact "where did this number come from?" discipline the workshop teaches.