

# Predictive AI for Engineers — Detailed Lesson Plan

---

**Text & Speech Analytics · four hands-on sessions · one six-hour day (09:30–15:30) · a no-code lab**

A one-day, four-session introduction to **predictive (classical) AI** through the two kinds of messy data engineers deal with every day: written **text** and spoken **speech**. This document gives the per-session timings, the demos, the case studies (with honest caveats and sources), the discussion prompts, and how the room's skeptics are handled.

The whole day deliberately covers the **predictive / analytical** half of AI — machine learning that *learns a pattern from past examples and predicts the next case*. Generative AI (ChatGPT and its relatives) is a different tool for a different job and is **out of scope here**, covered separately. Saying that out loud, early and often, is part of the plan.

## Who this is for, and the one goal

---

The room is **experienced mechanical, aeronautical, production and quality engineers** with no programming background — people who know their machines far better than we know them. Most arrive holding one of two beliefs: **"AI can do anything"** or **"AI is marketing nonsense."** Both are wrong, and the truth is more useful than either.

The single goal: every person leaves able to **reason about an AI claim** instead of believing or dismissing it — *what did it learn from, how was it tested on unseen cases, how confident is it, and where does it get things wrong?* We earn that by showing real, **honest, imperfect** results on data that looks like theirs. One overhyped claim and the skeptics tune out for the rest of the day, so credibility is the currency we spend carefully.

## Format — a no-code lab

---

- **No coding required.** Every code cell is pre-written. Participants only **run** cells (Shift + Enter) and read the output; in each session they change **one highlighted value** and re-run to see the result move.
- Runs in **Google Colab** in a browser — nothing to install. The notebooks are self-contained (built-in data, fixed seeds, so everyone sees the same numbers); one session also fetches a small public dataset live to prove the methods hold on real data.
- A rhythm repeated every session: **short framing → run a cell → "what just happened" → a plain-English vocabulary card → a knob to turn → recap.**
- **Mathematics is shown visually only** — a confusion-matrix grid, a similarity heatmap, a fitted line, a feature-importance bar, a spectrogram. Techniques are named for honesty (TF-IDF, logistic regression, FFT) but **no equations and no recall are expected.**

## The day at a glance

---

| # | Session (≈75 min)                                       | The one new idea                                                                                                              | The money-shot                                                                                                                     |
|---|---------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| 1 | <b>Text — sorting the words you already write</b>       | A computer can't read words, only count them: turn words into numbers (TF-IDF), then it's ordinary classification             | top clue-words per team — the model's reasoning in plain English                                                                   |
| 2 | <b>Text — predict, extract &amp; find repeat faults</b> | The same words-to-numbers front end answers many questions — a category, a number, a duplicate, a whole family, a part number | a similarity heatmap lighting up the same chronic fault, worded five different ways; a 2-D map of snags self-grouped into families |
| 3 | <b>Speech — voice → text → decision</b>                 | Speech-to-text is "audio in, text out"; then it's the text analytics from Sessions 1–2                                        | a low word-error transcript that still gets the one part number wrong                                                              |
| 4 | <b>Speech — predicting from the sound itself</b>        | A sound is already numbers; learn the pattern and a computer can <i>hear</i> a fault — no words at all                        | three fault types separating into clean clusters; the model leaning on the same whine and rumble a technician hears                |

*Day schedule (six hours, 09:30–15:30, breaks included): 09:30 Session 1 · 10:45 break · 11:00 Session 2 · 12:15 lunch · 12:45 Session 3 · 14:00 break · 14:15 Session 4 (closes with the day's wrap-up & Q&A) · 15:30 finish. Each session runs ≈75 minutes; the per-session plans below are the fuller menu — the blocks marked (bonus) and some discussion are trimmed to fit. All timings flexible.*

## Scope — four sessions, self-contained

The workshop is **four sessions, each built around one prepared Colab notebook**. The foundations of predictive AI (the five-step shape, the core vocabulary, the train/test idea) are **introduced inside Session 1 and reused all day** — there is no separate prerequisite. The same five-step ideas apply equally to **numeric sensor data** (classification, regression, predictive maintenance), but this workshop stays focused on text and speech.

## The single thread (put this on a wall poster)

Every session — text, speech, sound — is the *same five steps*. Master the shape once and the rest is variations:

**dataset → split (hide some) → train → predict → measure (honestly).**

Keep pointing back to it. The data changes shape (sentences, audio, a spectrum of pitches) but the picture never does. By Session 4 the room should be finishing the sentence for you.

## How we win the room (the doctrine behind every session)

- **Lead with credibility, not capability.** Open by naming the two mindsets and promising the honest middle. The first example of the day is the everyday **email spam filter** — classical text AI (bag-of-words, decades old, deployed at massive scale) that nobody calls magic.
- **Always show an imperfect result and open the confusion matrix.** For experienced engineers, visible errors are a *trust* signal; a flawless demo reads as a vendor trick. Every "it can" is paired with an honest "it can't."
- **The accuracy paradox, in their language.** If real failures are 1-in-1000, a model that always says "healthy" is 99.9% accurate and useless. That is why precision/recall and the confusion matrix — not a single accuracy number — are what an engineer should demand.
- **Predict-then-reveal, one knob.** Have the room commit to a guess out loud, change one highlighted value, re-run. Active prediction plus a single controlled variable produces the "aha" that passive watching never does.
- **Show, don't state.** Charts, never equations. Engineers read Pareto and control charts fluently; one formula on screen and a no-code audience checks out.
- **Respect their expertise.** Frame every text technique as *"reading the words you already write"* and every audio technique as *"a crude version of what your ear already does."* The model just learns their existing labelling habits; it rediscovers the cues they already trust; every extraction is auditable.  
**Auditability is the selling point.**
- **One aerospace + one factory example per technique**, so the skill is seen as general, not an aviation trick.
- **Be explicit about real vs synthetic vs simulated** — this audience will ask. Workshop demo data is synthetic and self-contained; named datasets (ASRS, CWRU, MIMII, MAFAULDA) are real; C-MAPSS is physics-simulated.
- **Close every demo with a prompt that pulls THEIR data into the room.** "What is the label column in *your* snag logs? Who assigns it today? Where would a wrong prediction actually cost you?" Concrete prompts get debate from skeptics; "any questions?" gets silence.

## Vocabulary, spread across the day

---

The basics are **not** front-loaded as a glossary. Each term is introduced the moment it is first needed, then reused relentlessly.

| Term                                | First met | In one plain line                                                                           |
|-------------------------------------|-----------|---------------------------------------------------------------------------------------------|
| dataset, row, column                | S1        | a table — one row per case, columns for what we recorded                                    |
| feature / label / label column      | S1        | features = the clues we get; label = the answer we want; the label column holds the answers |
| train/test split                    | S1        | learn from most, hide some to test honestly on unseen cases                                 |
| model, accuracy                     | S1        | the learned pattern; the fraction right on the hidden set                                   |
| bag-of-words, TF-IDF, stop words    | S1        | counting words; rare distinctive words count for more; ignore "the/and"                     |
| confusion matrix, the accuracy trap | S2        | a grid of what it got right and <i>how</i> it got things wrong                              |
| classification vs regression        | S2        | predict a bucket vs predict a number                                                        |
| cosine similarity                   | S2        | 0–1 score for how much two reports use the same distinctive words                           |
| embedding · clustering (k-means)    | S2        | a report as a point in space; k-means groups the points into families with no labels        |
| entity extraction (NER)             | S2        | pulling part numbers / components out of free text                                          |
| topic modelling (LDA)               | S2        | grouping co-occurring words into themes with no labels                                      |
| data leakage                        | S2        | giveaway clues sneaking into the test — a too-perfect score                                 |
| ASR, acoustic/language model        | S3        | speech-to-text; the two halves that turn sound into likely words                            |
| word error rate (WER)               | S3        | the speech version of accuracy — wrong + missing + extra words                              |
| keyword spotting                    | S3        | listening for a few fixed commands instead of every word                                    |
| sample rate, FFT, spectrum          | S4        | sound is numbers; the FFT splits it into the pitches it's made of                           |
| spectrogram, MFCC                   | S4        | a picture of pitch over time; the "professional" sound features                             |
| supervised vs anomaly detection     | S4        | learn from labelled faults vs learn "normal" and flag the surprising                        |
| orders (1×, 2×)                     | S4        | fault pitches read as multiples of shaft speed, not raw Hz                                  |

---

## Session 1 — Text: sorting the words you already write

**New idea:** words become numbers, then it's ordinary classification.

**By the end, participants can:** name the five steps of a predictive-AI project; use *dataset* / *feature* / *label* / *train-test split* / *accuracy* correctly; explain (in plain words) how text becomes numbers; and read a model's top clue-words to see it is not a black box.

| Time      | Segment                                   | What happens, and how it lands                                                                                                                                                                                                                                                                   |
|-----------|-------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 0:00–0:08 | <b>Frame the day</b>                      | Name the two mindsets. Promise the honest middle. The spam-filter anchor: you already trust classical text AI. State plainly: this is <i>predictive</i> , not generative.                                                                                                                        |
| 0:08–0:18 | <b>The five-step poster + first words</b> | Walk the spine: dataset → split → train → predict → measure. Introduce dataset/feature/label/label column as a plain spreadsheet idea.                                                                                                                                                           |
| 0:18–0:23 | <b>Into Colab</b>                         | Orient: Shift+Enter, run top-to-bottom, "Restart and run all" fixes almost everything.                                                                                                                                                                                                           |
| 0:23–0:53 | <b>Run the prepared notebook</b>          | Words → numbers (bag-of-words, stop words, TF-IDF) shown on two sentences first; vectorize the reports; train/test split; honest accuracy; <b>top clue-words per team</b> (oil/pressure/fluid → Hydraulic). Predict four fresh reports incl. one deliberately ambiguous — watch confidence drop. |
| 0:53–1:00 | <b>Turn the knob</b>                      | Change the train/test split ratio and re-run; discuss why the score wobbles on a small set.                                                                                                                                                                                                      |
| 1:00–1:12 | <b>Where this lives at work</b>           | Aerospace: auto-routing snags by ATA chapter; tagging safety narratives. Industrial: maintenance/IT ticket triage (an SVM hit ~85% on 3,632 HVAC maintenance requests); spam filtering.                                                                                                          |
| 1:12–1:22 | <b>Discussion</b>                         | "What is the label column in <i>your</i> logs? Who assigns it today, and how long does it take?"                                                                                                                                                                                                 |
| 1:22–1:30 | <b>Recap + bridge</b>                     | Five steps revisited; tease Session 2: sorting was the doorway — next we predict urgency, repair time, and catch repeat faults.                                                                                                                                                                  |

**Skeptic handling.** "*This is just if-then rules / statistics.*" — Yes, and that's the point: predictive AI is disciplined pattern-finding. The top-clue-words list proves it learned something human-readable, not magic. "*Why isn't it 100%?*" — On a small set the score is a rough indicator; a model that claims 100% on real data usually peeked at the answers (proven in Session 2).

**Real, citable anchors:** the standard text-classification benchmark reports ~83% (Naive Bayes) and ~91% (linear model) on a real public corpus — honest mid-80s/low-90s numbers to set expectations. NASA's Aviation Safety Reporting System holds 2.3M+ de-identified narratives (free, public-domain) that look just like HAL snag/incident logs; NASA itself calls them "soft data" it does not verify — an honest data-quality lesson in one sentence.

---

## Session 2 — Text: predict, extract & find recurring faults

**New idea:** the same words-to-numbers front end answers many questions — just swap the final step.

**By the end, participants can:** tell classification from regression and know which question each answers; read a confusion matrix for *what* a model gets wrong; explain how duplicate/recurring faults are found with no labels and how clustering groups all the snags into families (embeddings as points in space); see how part numbers are lifted out of prose and audited; and recognise data leakage (the too-perfect score).

| Time          | Segment                                              | What happens, and how it lands                                                                                                                                                                                                                                                                                                                                                                                                                 |
|---------------|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 0:00–<br>0:06 | <b>Recall + preview</b>                              | One-line recap of "words → numbers." Today: predict a category, predict a number, find repeats, pull out parts.                                                                                                                                                                                                                                                                                                                                |
| 0:06–<br>0:12 | <b>The snag log</b>                                  | A 300-row synthetic MRO log: free-text report, system, urgency, downtime hours. Note the last is a <i>number</i> .                                                                                                                                                                                                                                                                                                                             |
| 0:12–<br>0:28 | <b>Predict a CATEGORY (urgency)</b>                  | TF-IDF + logistic regression → ~81% on unseen snags. Open the <b>confusion matrix</b> : mistakes are mostly Low↔Medium; it caught 10/15 High snags and slipped High→Low only once — <i>rare, but it happened</i> , which is exactly why the High bucket keeps a human's eyes. Fold in the <b>accuracy trap</b> : a model that always says "routine" can score high yet catch zero urgent snags.                                                |
| 0:28–<br>0:42 | <b>Predict a NUMBER (downtime)</b>                   | Same TF-IDF in, swap to a regressor → off by ~1.5 h on ~6 h jobs ( $R^2 \approx 0.6$ ). The predicted-vs-actual scatter makes "it learned the relationship" visible — and honest about the big-job misses. <i>This is the classification-vs-regression "aha"</i> .                                                                                                                                                                             |
| 0:42–<br>0:54 | <b>Find RECURRING faults</b>                         | TF-IDF + cosine similarity on snags worded differently. The heatmap lights up two blocks — the same chronic fault, de-duplicated, no labels.                                                                                                                                                                                                                                                                                                   |
| 0:54–<br>1:04 | <b>Group into families (embeddings + clustering)</b> | Embed each snag as a point in space, then <b>k-means</b> sorts <i>all</i> of them into 5 families on a 2-D map — no labels. A few families pop out cleanly (Mechanical, Electrical, Avionics) but a catch-all swallows the rest, so agreement with the real systems is honestly <b>low (~0.22)</b> — k-means groups by language, not your org chart; you read the top words and name each. The unsupervised companion to the duplicate finder. |
| 1:04–<br>1:12 | <b>Pull out PART NUMBERS</b>                         | A few readable rules extract part numbers into an auditable table, then a "most-mentioned parts" chart — free text becomes a reliability/spares signal. Every hit is checkable.                                                                                                                                                                                                                                                                |
| 1:12–<br>1:22 | <b>Bonus: themes (LDA) + the honest-test trap</b>    | One pre-baked topic view (some topics map to systems, some to severity language — <i>you</i> name them). Then the same pipeline on a real public corpus: ~84% honestly vs ~95% when giveaway headers leak in. <i>How was it tested, and could it have peeked?</i>                                                                                                                                                                              |
| 1:22–<br>1:30 | <b>Discussion + recap</b>                            | "Where would a wrong prediction be expensive vs cheap for you? Do your historical logs exist <i>labelled</i> ?"                                                                                                                                                                                                                                                                                                                                |

**Skeptic handling.** *"Can it replace our engineers?"* — No: it hands you a prioritised list and a reason; a human decides. The confusion matrix and the High→Low slip make the limits measurable, which is exactly why it's trustworthy.

**Real, citable anchors (with honest caveats):**

- **Chronic/recurring-defect detection** is a live commercial product in aviation MRO. Its headline figures — used by >25% of the world fleet, ~40% of defects unflagged due to inconsistent coding, 90%+ code prediction — are **vendor-claimed, not independently audited**. Lead with the capability (de-duplicating differently-worded faults), not the percentages.
- **Part-number extraction**: a US defence study lifted part-number extraction accuracy (F1) from ~3% (plain pattern-matching) to ~95% (a trained extraction model) — an honest *mixed* result (rigid codes were easy for both).
- **Defence precedent**: a national defence organisation mined two years of military-aircraft maintenance and pilot free-text to surface "lead indicators" of defects — the closest published analogue to HAL's situation.
- **Topic modelling, honestly**: a NASA-data study recovered >70% of hand-curated reports on an emerging theme via topic modelling — *but* a separate study concluded off-the-shelf topic tools were "insufficient for sense-making" on their own. A useful map, not a verdict.
- **Industrial transfer**: national consumer-complaint databases (~1.9M free-text records, public) are the warranty-mining analogue.

*Indian context to mention lightly*: after a 2025 accident, India's DGCA is reported to be tightening repetitive-defect reporting (a defect recurring three times within 15 flight cycles) — **press-reported/draft**, so cite it as motivation, not settled regulation.

---

## Session 3 — Speech: voice → text → decision

---

**New idea**: speech-to-text is "audio in, text out"; then it's the text analytics you already know.

**By the end, participants can**: describe ASR as the same predict-from-examples idea with sound as the input; read WER as the speech version of accuracy and explain why a low WER can still be unsafe; and say why a small fixed vocabulary (keyword spotting) is more reliable than full transcription.



| Time          | Segment                                | What happens, and how it lands                                                                                                                                                                                                                                                                                       |
|---------------|----------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 0:00–<br>0:08 | <b>Frame</b>                           | The last messy data type: spoken words. The pipeline: speak → transcribe → decide. Draw the classical→modern line (HMM-GMM → deep models, ~2012) explicitly so nobody thinks "this is ChatGPT" — both are <i>predictive</i> , not generative.                                                                        |
| 0:08–<br>0:30 | <b>Run the speech pipeline</b>         | Manufacture a spoken report (text-to-speech, no mic needed) → transcribe it → route it with the Session-1 classifier, with a confidence bar. ( <i>First cell installs two tools, ~1 min — the only setup of the day.</i> )                                                                                           |
| 0:30–<br>0:45 | <b>Measure it honestly (WER)</b>       | A self-contained WER demo: a noisy transcript scores a worse 15% yet stays harmless; a "cleaner" 8% transcript quietly turns one part number into another — the one word that matters, now wrong. <b>Low error rate ≠ safe.</b> Rule: a human verifies numbers and serials.                                          |
| 0:45–<br>0:55 | <b>Keyword spotting</b>                | Listening for a 4-command vocabulary inside messy transcripts — robust to stutters and filler, and it leaves the no-command line alone. Why hands-free shop-floor tools use wake-words, not dictation.                                                                                                               |
| 0:55–<br>1:08 | <b>Where it works, where it breaks</b> | Hard by default (accents, jargon, radio noise), strong with domain data: a European air-traffic-control speech project reached >95% (controllers) / >90% (pilots) word recognition <i>after</i> domain training, ~97% on callsigns when fused with radar. Open ATC corpora exist (clean and noisy) for exactly this. |
| 1:08–<br>1:20 | <b>Discussion</b>                      | "Where do people <i>speak</i> information that never gets written down — handovers, inspections, calls? What's the one word in your world that must never be mis-heard?"                                                                                                                                             |
| 1:20–<br>1:30 | <b>Recap + bridge</b>                  | Speech → text → decision, honestly measured. Tease Session 4: next we throw the words away and predict from the <i>sound itself</i> .                                                                                                                                                                                |

**Skeptic handling.** *"ASR is hype — it never gets my accent."* — Agree out loud, then show the WER demo and the domain-adaptation counter-evidence: out of the box it struggles; trained on your domain it gets strong. Honesty about the limit is what makes the capability believable.

---

## Session 4 — Speech: predicting from the sound itself

**New idea:** a sound is already numbers; learn the pattern and a computer can *hear* a fault — no words.

**By the end, participants can:** explain that a sound is just numbers and that a few features (energy, low/high bands, dominant pitch) capture what the ear hears; tell supervised classification from unsupervised anomaly detection and which fits the realistic "faults are rare" case; and read fault pitches as orders of shaft speed.



| Time      | Segment                                         | What happens, and how it lands                                                                                                                                                                                                                                                                                                                           |
|-----------|-------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 0:00–0:08 | <b>The technician's ear</b>                     | A seasoned tech walks past and says "that bearing's going" — from the sound. Can a computer learn that? Same idea, last time: numbers → pattern → judge.                                                                                                                                                                                                 |
| 0:08–0:22 | <b>Sound is numbers; listen, then featurise</b> | Sample rate → a clip is a list of numbers. <i>Play</i> healthy / bearing / imbalance clips. The FFT splits a sound into its pitches (a prism); five features capture loudness, low rumble, high whine, dominant pitch, brightness.                                                                                                                       |
| 0:22–0:34 | <b>See the pattern, then classify</b>           | Scatter of low-band vs high-band energy → three clean clusters. Train/test split → ~88% on unseen clips. <b>Feature importance</b> shows the model leaning on the whine and the rumble — the very cues the ear uses.                                                                                                                                     |
| 0:34–0:44 | <b>Diagnose a brand-new machine</b>             | Play a fresh clip, get a fault call + confidence. Turn the knob: try other faults / seeds; widen the bearing range and watch accuracy soften — honest.                                                                                                                                                                                                   |
| 0:44–0:58 | <b>When you have no fault examples</b>          | The realistic aerospace case: lots of healthy sound, few faults. Train an anomaly detector on <b>healthy only</b> (AUC ≈ 0.94); the histogram shows healthy left, faults right of the alarm line. Move the line: ~5% false alarms catches ~96% of imbalance and ~ $\frac{2}{3}$ of bearing faults — <i>no free lunch</i> , and the trade-off is visible. |
| 0:58–1:10 | <b>Real datasets + orders</b>                   | Name the real, benchmarked field: CWRU bearings, MAFAULDA (with a mic channel), MIMII / the DCASE challenge (honest AUCs ~0.54–0.96). Read pitches as <b>orders</b> (imbalance 1×, misalignment 2×), not raw Hz.                                                                                                                                         |
| 1:10–1:20 | <b>Where it lives in aerospace</b>              | Helicopter HUMS detects ~70% of the failure modes it monitors (strong, explicitly imperfect); engine health monitoring (one OEM monitors 8,000+ aircraft); acoustic-emission NDT for cracks — same recipe: signal → features → classify/flag.                                                                                                            |
| 1:20–1:30 | <b>The day in one picture + close</b>           | Numbers, text, speech, sound — it was always the same five steps. The "neither magic nor nonsense" close: you can now <i>question</i> an AI claim. Hand off to the generative-AI sessions to come.                                                                                                                                                       |

**Skeptic handling.** *"But we don't have failure labels."* — Exactly why the anomaly-detection demo matters: learn normal, flag the surprising, no fault examples needed. *"Why not 100%?"* — public benchmark scores run from ~54% to ~96%; a flawless demo is the thing to distrust.

---

## Case studies (and how to cite them honestly)

| Case                                                                      | What it really is                                                                                                   | How to say it                                                                                                                                                                             |
|---------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Airbus Skywise — A330neo bleed valve (EASA emergency AD, Aug 2022)</b> | Cross-fleet data analysis flagged a software bug over-stressing a valve pin; caught <i>before</i> any accident      | The best honest safety story for the room — predictive analytics as an engineering early-warning, not cost-savings hype. (Skywise now connects 12,000+ aircraft — quote the live figure.) |
| <b>Delta TechOps predictive maintenance</b>                               | Maintenance-caused cancellations fell from 5,600+ (2010) to ~55 (2018), >95% success on pending-failure predictions | Powerful — but <b>self-reported</b> , not independently audited. Model the "how was this measured?" discipline.                                                                           |
| <b>Aviation chronic-defect NLP</b>                                        | Reads maintenance logbook free-text to recognise the same chronic defect described differently                      | <b>Vendor-claimed</b> figures; lead with the capability, flag the numbers as unaudited.                                                                                                   |
| <b>Air-traffic-control speech recognition</b>                             | Domain-adapted ASR for ATC                                                                                          | >95% / >90% word recognition <i>after</i> domain training; ~97% callsigns with radar context. Frame as "hard by default, strong with domain data."                                        |
| <b>Machine-sound anomaly benchmark (DCASE / MIMII)</b>                    | Public benchmark for detecting machine faults from sound using only normal data                                     | The anti-hype anchor: honest AUCs ~0.54–0.96 across machine types.                                                                                                                        |
| <b>Engine health monitoring</b>                                           | Streaming engine parameters to forecast component replacement, underpinning "power-by-the-hour" service             | One OEM reports monitoring 8,000+ aircraft. Shows the business model that forces honest models.                                                                                           |
| <b>Indigenous fighter-fleet predictive maintenance (India, 2026)</b>      | A prognostic maintenance programme for a ~270-strong fleet, with AI and real-time sensor monitoring                 | "This is happening here" context — but it is <b>sensor/prognostics</b> , not text/audio; cite as context only.                                                                            |

On HAL itself: there is **no verified public evidence** of HAL running predictive ML in production. Frame everything as "**applicable to HAL's operations**," not "HAL already does this." HAL's strongest latent asset is its own historical snag/overhaul logs — *if* they are clean and labelled.

## Public datasets (Colab-usable or citable)

| Dataset         | What                                                                  | Use                                       | Source                                    |
|-----------------|-----------------------------------------------------------------------|-------------------------------------------|-------------------------------------------|
| NASA ASRS       | 2.3M+ de-identified aviation safety narratives (free text + synopsis) | S1–S2 text routing / topics               | asrs.arc.nasa.gov (public domain)         |
| MaintNet        | 6,169 real aviation maintenance logbook entries + tools               | S1–S2 "cleaning beats the model"          | aclanthology.org/2020.coling-demos.2      |
| FAA SDRS        | ~1.7M service-difficulty records, free-text + structured              | S1–S2 trend/routing                       | faa.gov/av-info/download_SDR              |
| 20 Newsgroups   | ~18k real documents, one-line fetch                                   | S2 real-data proof + the header-trap      | built into scikit-learn                   |
| ATCOSIM / ATCO2 | Clean / noisy air-traffic-control speech corpora                      | S3 realism                                | spsc.tugraz.at · arxiv.org/abs/2211.04054 |
| Speech Commands | 1-sec spoken-word clips                                               | S3 keyword spotting                       | arxiv.org/abs/1804.03209                  |
| CWRU Bearing    | The bearing-fault benchmark                                           | S4 "real, benchmarked field"              | engineering.case.edu/bearingdatacenter    |
| MIMII           | Real valve/pump/fan/slide-rail sound + factory noise                  | S4 anomaly detection                      | zenodo.org/records/3384388                |
| MAFAULDA        | 1,951 runs: imbalance/misalignment/bearing, incl. a mic channel       | S4 maps to hand-diagnosed faults          | www02.smt.ufrj.br/~offshore/mfs           |
| NASA C-MAPSS    | <i>Simulated</i> turbofan run-to-failure (remaining useful life)      | S4 regression bridge (label it simulated) | data.nasa.gov                             |

*Heavy downloads (CWRU, MIMII, MAFAULDA) stay out of the live lab — they are named and linked for credibility; the synthetic demos are the floor, not the ceiling.*

## Foundations & optional material

The five-step foundations — *dataset, feature, label, label column, train/test split, accuracy, the confusion matrix* — are introduced inside **Session 1** and reused all day; there is no separate prerequisite. The same ideas apply just as well to **numeric sensor data** (classification, regression, predictive maintenance with feature importance and anomaly detection), but the workshop stays focused on text and speech.

## Lab logistics & pre-flight

- A computer lab with **internet**, one machine per participant (pairs are fine), able to reach Google Colab.
- A **Google account** per participant (personal works), or the upload fallback.
- A projector for the framing slides and the money-shot charts.
- **Pre-flight (day before):** run **Session 3** end-to-end once — its first cell installs the speech tools (~1 minute) and downloads a small model on first transcription; confirm the network allows it. Everything else needs no install.
- **The one reset that fixes almost everything:** *Runtime* → *Restart and run all*. Say it often.

## Sources & further reading (selected, verified)

| Topic                                                                                 | Source                                                                                                                              |
|---------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| Spam filtering as classical text AI                                                   | <a href="https://en.wikipedia.org/wiki/Naive_Bayes_spam_filtering">en.wikipedia.org/wiki/Naive_Bayes_spam_filtering</a>             |
| Honest text-classifier accuracy + the header trap                                     | <a href="https://scikit-learn.org">scikit-learn.org</a> (20 Newsgroups tutorial & dataset notes)                                    |
| Work-order text classification (SVM ~85% / 3,632 HVAC)                                | ASCE J. Arch. Eng. (2022), doi 10.1061/(ASCE)AE.1943-5568.0000522                                                                   |
| Aviation Safety Reporting System & "soft data" caveat                                 | <a href="https://asrs.arc.nasa.gov">asrs.arc.nasa.gov</a>                                                                           |
| Part-number extraction (3%→95%)                                                       | <a href="https://apps.dtic.mil/sti/trecms/pdf/AD1157022.pdf">apps.dtic.mil/sti/trecms/pdf/AD1157022.pdf</a>                         |
| Defence maintenance-text NLP (ICAS 2024)                                              | <a href="https://icas.org/icas_archive/icas2024">icas.org/icas_archive/icas2024</a> (#0083)                                         |
| Topic modelling on safety narratives (>70% recovery; "insufficient for sense-making") | <a href="https://ntrs.nasa.gov/citations/20210011508">ntrs.nasa.gov/citations/20210011508</a> ; /20205003750                        |
| Word Error Rate definition                                                            | <a href="https://en.wikipedia.org/wiki/Word_error_rate">en.wikipedia.org/wiki/Word_error_rate</a>                                   |
| Air-traffic-control speech recognition (HAAWAI)                                       | <a href="https://cordis.europa.eu/article/id/442201">cordis.europa.eu/article/id/442201</a>                                         |
| Deep models replace GMM (~2012)                                                       | <a href="https://cs.toronto.edu/~hinton/absps/DNN-2012-proof.pdf">cs.toronto.edu/~hinton/absps/DNN-2012-proof.pdf</a>               |
| Machine-sound anomaly benchmark (DCASE 2020, AUC ~0.54–0.96)                          | <a href="https://ar5iv.labs.arxiv.org/abs/2006.05822">ar5iv.labs.arxiv.org/abs/2006.05822</a>                                       |
| Helicopter HUMS (~70% of monitored failure modes)                                     | <a href="https://skybrary.aero/articles/vibration-health-monitoring-vhm">skybrary.aero/articles/vibration-health-monitoring-vhm</a> |
| Skywise A330neo early warning (EASA AD 2022)                                          | <a href="https://theaircurrent.com">theaircurrent.com</a> (Skywise A330neo)                                                         |
| Engine health monitoring (power-by-the-hour)                                          | <a href="https://rolls-royce.com">rolls-royce.com</a> (civil aerospace / EHM)                                                       |

*All quantitative claims above were fact-checked; figures that are vendor-claimed, self-reported, simulated, or press-reported/draft are flagged as such in the text — modelling the exact "where did this number come from?" discipline the workshop teaches.*